Asha Vidyadharan
[Department], [Institution]
[City], [Country]

June 27, 2026

The Editor-in-Chief
[Target Journal Name]

Dear Editor,

We are pleased to submit our manuscript, **"DarkTrace: An Explainable and Forensically Verifiable Multilingual Dark Web Threat Intelligence Framework with Severity Scoring and Evidence Integrity Validation,"** for consideration as a research article in *[Target Journal Name]*.

Cyber threat intelligence increasingly originates from the dark web, yet existing automated pipelines share three limitations that, critically, co-occur in real investigations: they are predominantly English-centric, they expose opaque "black-box" classifications that analysts cannot defend operationally or legally, and they do not preserve the evidentiary integrity required for their outputs to be admissible as digital evidence. To our knowledge, no prior system addresses these jointly.

Our central contribution is the *integration*. DarkTrace unifies, in one investigator-ready pipeline: (i) multilingual threat classification spanning English and non-English content; (ii) a hybrid, explainable, calibrated severity-scoring model with SHAP/LIME/attention attributions; (iii) STIX 2.1-aligned actor/entity extraction; and (iv) an end-to-end forensic evidence-integrity layer (SHA-256 hashing, trusted timestamping, append-only chain-of-custody) that makes forensic soundness a cross-cutting property of the whole pipeline rather than an after-the-fact step.

We believe the work is well suited to *[Target Journal Name]* because it advances both the methodology and the practice of trustworthy, legally defensible AI-assisted threat intelligence. Key empirical findings, obtained on sanitized, archived, and public corpora and reported with confidence intervals and significance tests throughout, include:

- dark-web text classification at 0.919 macro-F1, with the model differences statistically validated (McNemar $p = 4.2 \times 10^{-4}$);
- multilingual detection across six languages plus Hindi/Arabic cross-domain, with cross-lingual transfer characterized and small-sample languages bounded by exact Wilson intervals;
- an explainable severity scorer that *statistically outperforms* a black-box baseline on thresholded macro-F1 (+0.022, 95% CI $[0.011, 0.034]$, $p < 0.001$, Holm-corrected) while adding interpretability and improved calibration;
- 100% tamper detection in an adversarial evidence-integrity test at $\sim$9.9k items/s, with valid STIX 2.1 export; and
- an integration ablation showing that only the full pipeline maximizes the actionability of the emitted intelligence.

We confirm that this manuscript is original, has not been published previously, and is not under consideration elsewhere. The work raises no undisclosed ethical concerns: all data are sanitized, archived, public, or synthetic; no illegal collection was performed; personal and victim data are minimized and never redistributed; and our ethics and legal protocol is described in the manuscript. All authors have approved the submission and declare no competing interests. Code, derived data, and result artifacts are released to support reproducibility.

We would be glad to suggest qualified reviewers and to address any questions. Thank you for your time and consideration.

Sincerely,

Asha Vidyadharan
`codeasasha@gmail.com`